UFMG UNIVERSIDADE FEDERAL DE MINAS GERAIS

Ranking Models

Introduction

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br

Information retrieval

> *Information retrieval is a field concerned with the structure, analysis, organization, storage, searching, and retrieval of information.*
>
> Gerard Salton, 1968

Retrieval tasks

Retrieval tasks

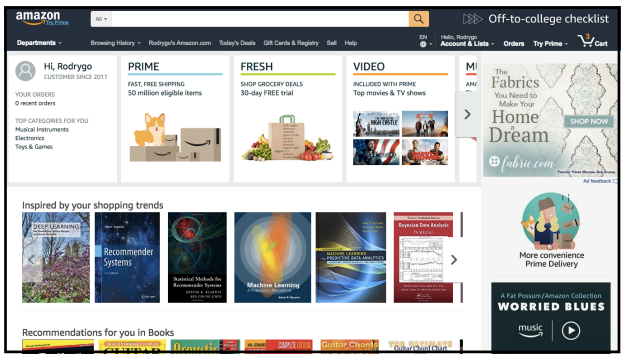

Retrieval tasks

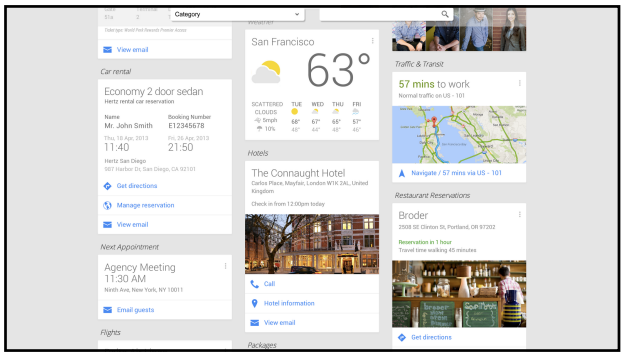

Our focus: search

A lot of people ..... *60k queries per second*
From a lot of places ..... *whole planet (and beyond?)*
Using a lot of a devices ..... *smart-you-name-it*
Looking for a lot of info ..... *10^{11} documents*

What does a search engine do?

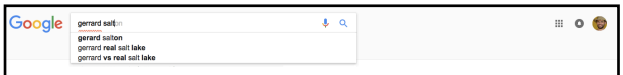

Spelling correction

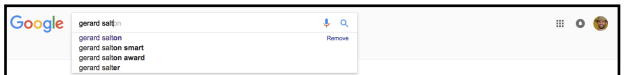

Query autocompletion
Query suggestions

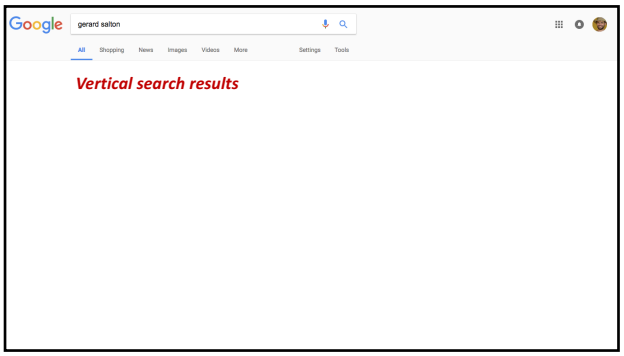

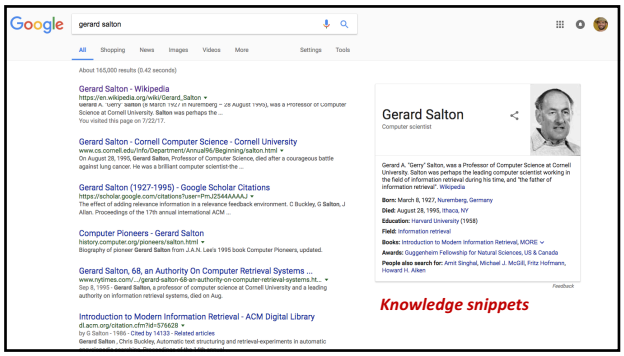

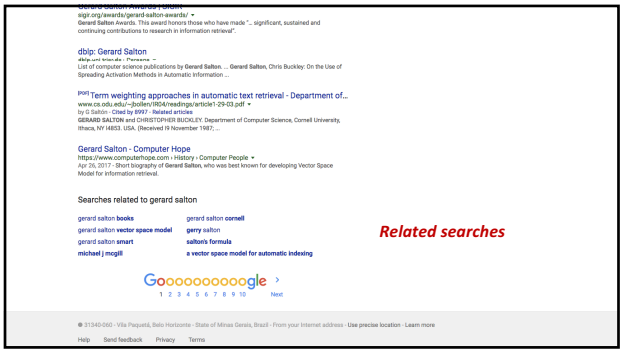

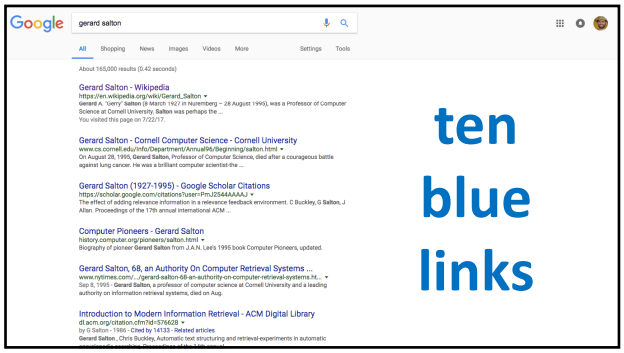

The ranking problem

Given

- Some evidence of the user's need

Produce

- A list of matching information items
- In decreasing order of relevance

The ranking problem

Given

- Some ~~evidence of the user's need~~ ***query***

Produce

- A list of matching ~~information items~~ ***documents***
- In decreasing order of relevance

1) What documents do we show?

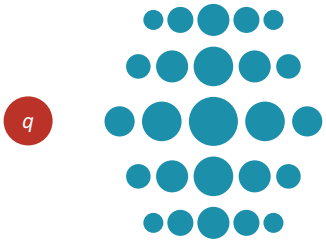

2) What order do we show them in?

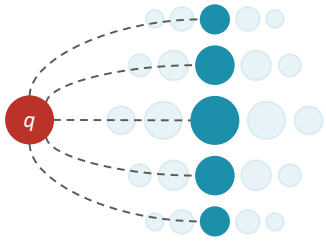

2) What order do we show them in?

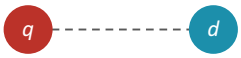

$f(q, d)$

Isn't it a solved problem?

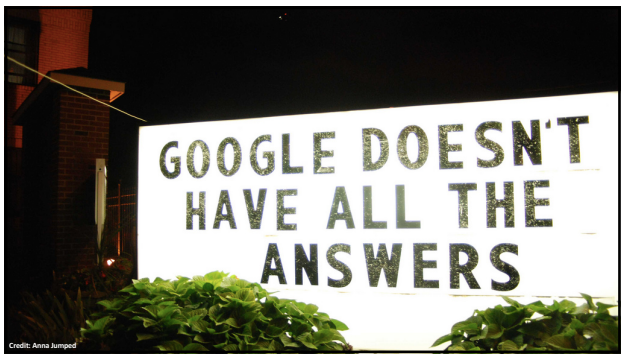

Efficiency

Efficiency is about doing something (good or bad) in an optimal way (i.e., faster or with fewer resources)

Key performance indicators

- *Query latency:* searching billions of documents
- *Query throughput:* serving thousands of users
- *Document latency:* serving freshly published content

Effectiveness

Effectiveness is about doing the right thing; it's about finding documents that are relevant to the user

Relevance is influenced by many factors
- Topical relevance vs. user relevance
- Task, context, novelty, style

Ranking models define ***a view of*** relevance

Pursuing relevance

Exact matching of words is not enough
- Queries are ambiguous; documents are ambiguous
- Many different ways to write the same thing

Even perfect matching is not enough
- Must counter adversarial content
- Must infer quality beyond content

Assessing relevance

Relevance is a user's prerogative
- We can observe changes in user behavior
- Or directly ask the user how we're doing

Evaluation is an empirical science
- It must be scientifically rigorous
- It must be economically viable

What do search engineers do?

The ranking problem

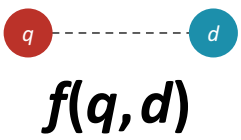

$f(q, d)$

Ranking pipeline

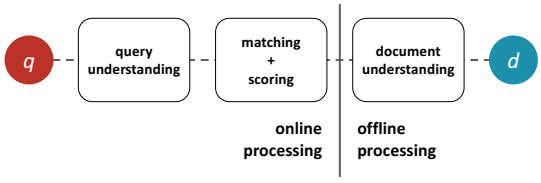

(Continuous) offline processing

Document acquisition

Document understanding

Document indexing

Document acquisition

The Web is huge

- Trillions of known URLs, billions fetched

The Web is constantly evolving

- Updates, additions, deletions

Efficient crawling is key

- Must aim for coverage, but also freshness

Document understanding

Documents come in many flavors

- Web pages, email, books, news stories, scholarly papers, text messages, Word™, Powerpoint™, PDF, forum postings, patents, IM sessions, etc.

Document understanding

Documents have mostly textual content

- Some of it static, some dynamically rendered

And often some structure

- Title, body, url, anchor text for web pages
- Subject, sender, destination for email

Document understanding

Documents carry meaning

- Term-based matching as a first approximation
- Several techniques to leverage semantics

Documents vary in quality

- *Genuinely:* accessibility, readability, authority, depth
- *Maliciously:* content farms, link farms

Document indexing

Efficient retrieval through indexes

- Like the index of a book
  - For each word, a list of documents it appears on
- Broken up into shards of millions of documents
  - 1000s of shards for the web index
- Plus per-document metadata

Online processing

Query understanding

Matching and scoring

Post-processing

Query understanding

Keyword queries are often poor descriptions of the user's actual information need

- Interaction and context also matter

Query understanding techniques can help

- Query segmentation, query scoping
- Query relaxation, query expansion

Query understanding

Query scoping through semantic annotation

- [**san jose** convention center]
- [**matt cutts**]

Query expansion through acronym expansion

- [**gm** trucks] → [**general motors** trucks]
- [**gm** corn] → [**genetically modified** corn]

Matching and scoring

Send the query to all the shards

Each shard

- Finds matching documents
- Scores each query-document pair
- Sends back the top *n* documents

Combine all the top documents and sort by score

Post-retrieval adjustments

Host clustering, sitelinks

Near-duplicate removal, diversification

Spam demotions, copyright takedowns

Course goals

Provide an in-depth account of ranking models and evaluation methods for information retrieval

Provide an exploration of recent advances and current research directions in the field

Course scope

Focus on ranking
- Query-dependent ranking
- Query-independent ranking
- Machine-learned ranking

Focus on evaluation
- Offline, online, counterfactual

Out-of-scope

We have dedicated courses for:
- Information retrieval
- Recommender systems
- Machine learning
- Data mining

Course materials: textbooks

Search Engines: Information Retrieval in Practice
by B. Croft, D. Metzler, and T. Strohman

Introduction to Information Retrieval
by C. Manning, P. Raghavan, and H. Schütze

Modern Information Retrieval
by R. Baeza-Yates and B. Ribeiro-Neto

Course materials: textbooks

Information Retrieval: Implementing and Evaluating Search Engines
by S. Büttcher, C. Clarke, and G. Cormack

Text Data Management: A Practical Introduction to Information Retrieval and Text Mining
by C. Zhai and S. Massung

Course materials: surveys

Foundations and Trends in Information Retrieval
by several authors

Morgan & Claypool Synthesis Lectures on Information Concepts, Retrieval, and Services
by several authors

Other relevant material

General background
- Algorithms and data structures
- Basic statistics
- Basic linear algebra

Advanced readings
- Google Scholar is your friend

Course grading (tentative)

Exams: 40%

Project: 40%

Assignments: 20%

Course website

http://homepages.dcc.ufmg.br/~rodrygo

References

Google Search Statistics
Internet Live Stats, 2017

How Google Works: A Google Ranking Engineer's Story
Haahr, SMX West 2016

Ten blue links on Mars
Clarke et al., WWW 2017

Writing assignment #0

Fill in a short questionnaire describing your past experience and expectations related to the course

- **Due Tue, Aug 15 @ 23:55 via Moodle**

UFMG UNIVERSIDADE FEDERAL DE MINAS GERAIS

Coming next…

Search Architecture

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br